To churn or not to churn?

The KKBOX Kaggle Challenge

100% Python

DATA MINING I PROJECT
TEAM 9
HWS 2017

KKBOX and the meaning of Churn

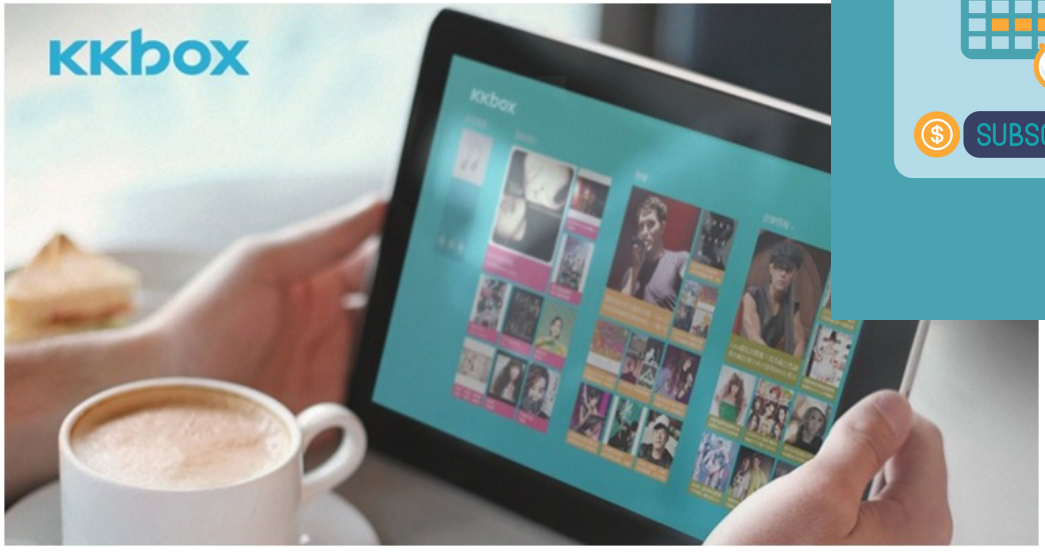

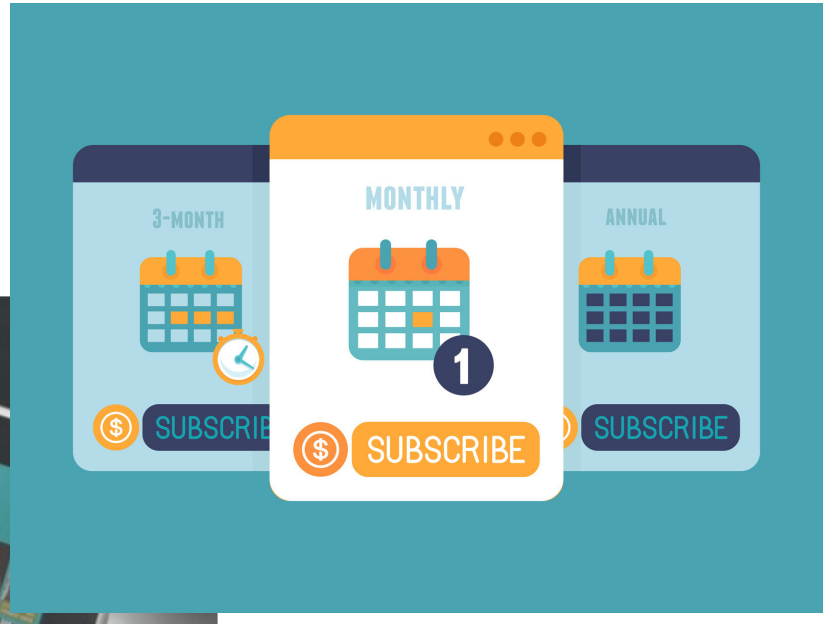

The DATA

members

user logs

transactions

train

Data Preparation

Reduce height : Data selection and aggregation

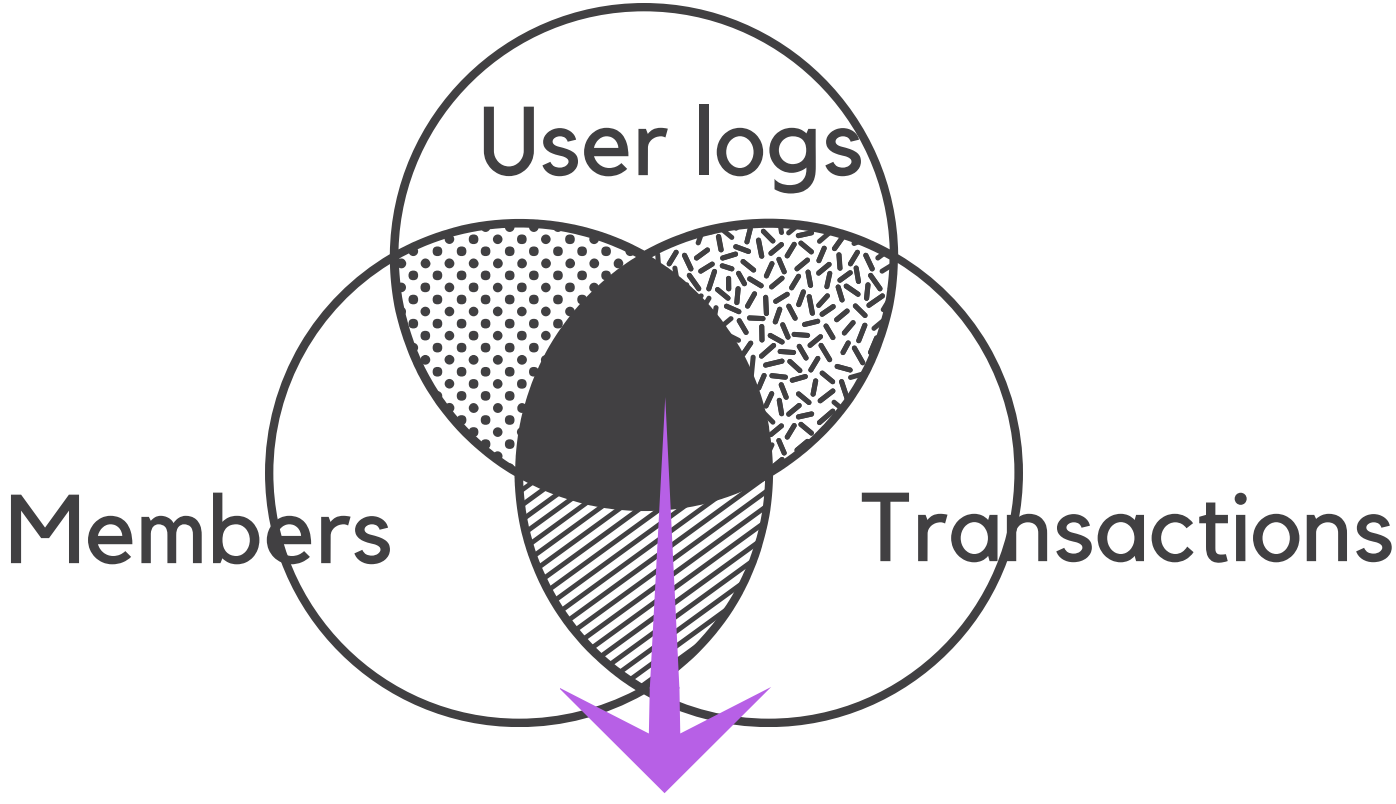

6 months view

Data Cleaning and Transformation

Generating new attributes:
- total_churn based on transactions
- monthly aggregated usage activities

Removing abnormal values:
- age up to 2000?
- days below 0?

Treating missing and null values:
- categorical: replace with most common
- numeric: replace with mean

Split and standardize train set

RobustScaler

Undersample

Apply ANN and DT

Handle outliers with isolation forest

StandardScaler

Undersample

Apply ANN and DT

Balance train dataset and evaluate

RobustScaler

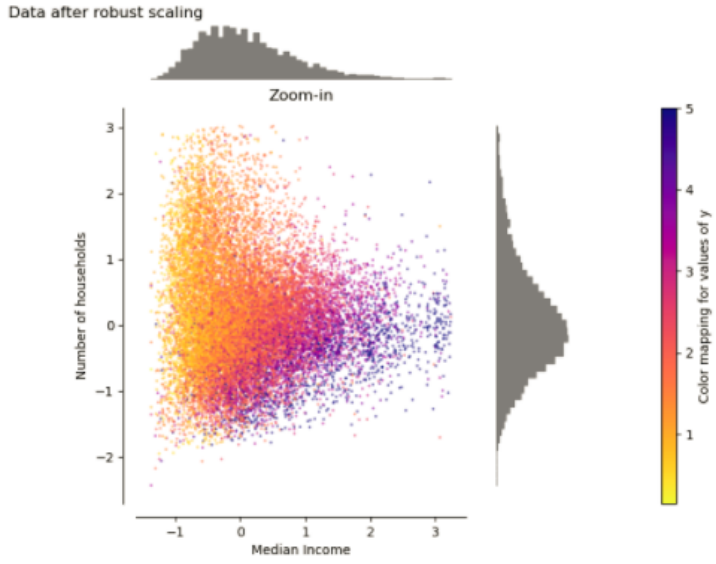

ANN: 90.82
DT: 88.59

StandardScaler

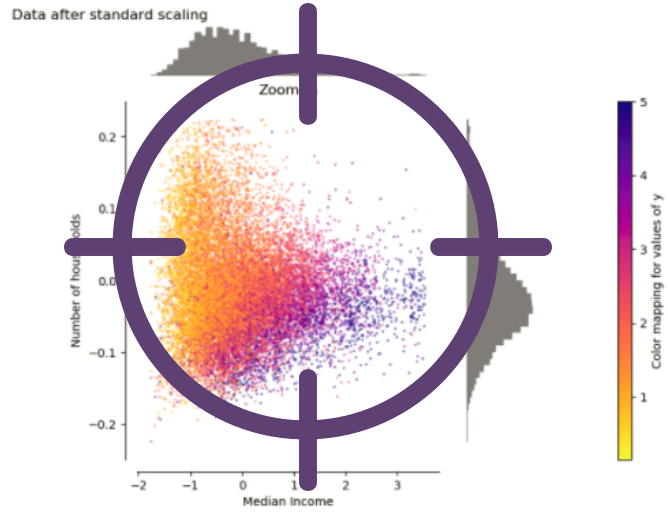

ANN: 91.89
DT: 89.66

Reduce width: Feature selection

154 features

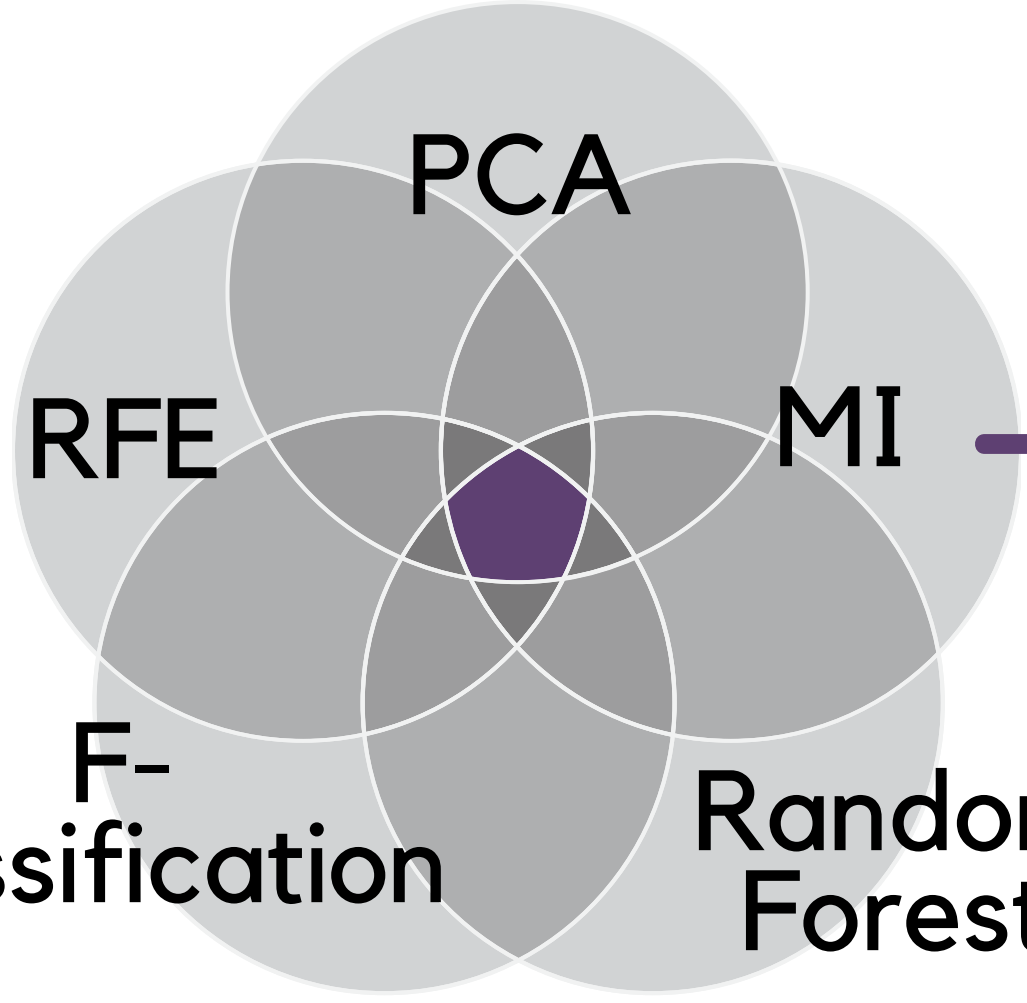

59 features

Reduce width: Correlation Matrix

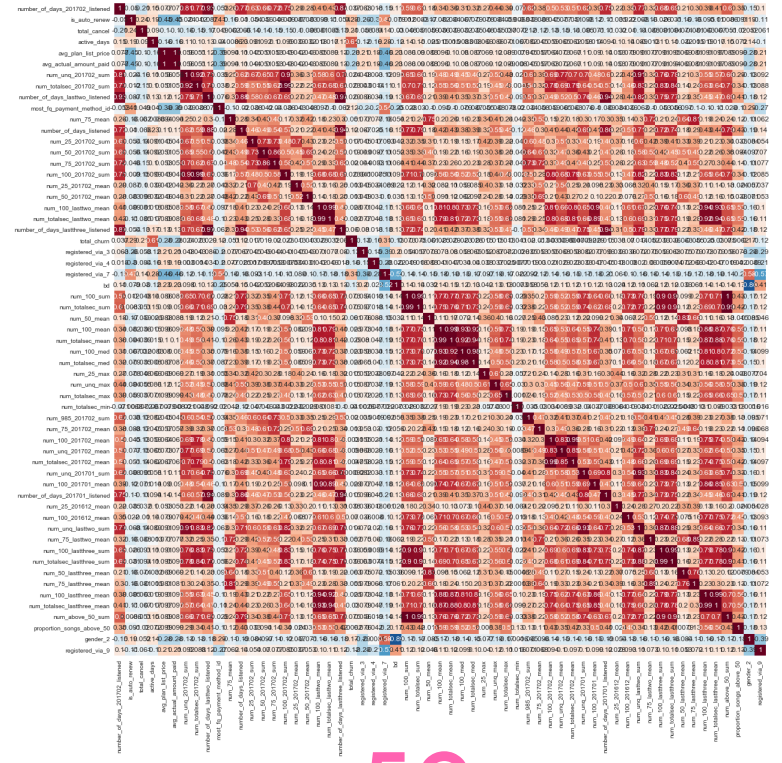

59

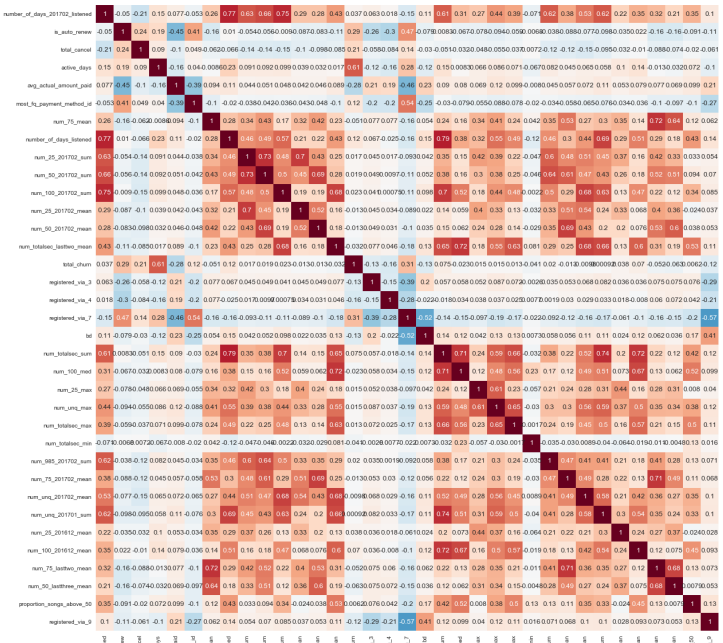

35

Data Mining

52 Models

Logistic Regression
SVM
K-NN
Nearest Centroid
Naive Bayes
Decision Tree
ANN

In Search for the best model...

Approach: start with a baseline model for each algorithm, fine-tune the parameters locally and then compare the chosen best models globally .

- K-NN : optimal number for K?
- DT : combination of parameters, pure nodes removal?
- ANN : optimiser function, features, hidden layers/units?
- Logistic Regression : optimal features subset, C value?
- SVM : kernel, optimal gamma, C ?
- Naive Bayes/ Nearest Centroid : optimal features subset?

Grid Search

Evaluation Metrics

- Not accuracy: count of correct predictions

- Recall , Precision and F1-score

- ROC curve and AUC

- **Log-loss** : account for the uncertainty of your prediction

$$logloss = -\frac{1}{N}\sum_{i=1}^{N}\left(y_i \log(p_i) + (1-y_i)\log(1-p_i)\right)$$

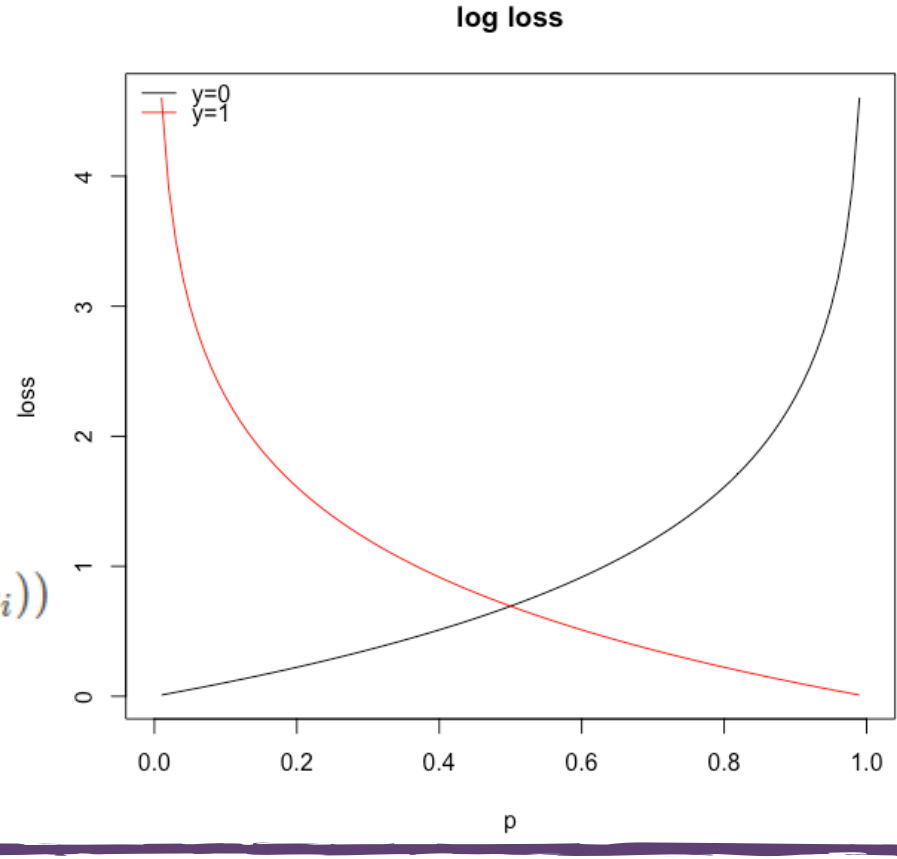

Performance Evaluation

Evaluation	K-NN	NC	LogR	DT	ANN	NB	SVM
AUC	0,96	-	0,95	0,97	0,97	0,94	0,95
Log Loss	0,82	-	0,27	0,20	0,21	0,63	0,27
Precision-avg/total	0,95	0,95	0,95	0,95	0,96	0,95	0,95
Recall-avg/total	0,89	0,89	0,89	0,90	0,91	0,84	0,89
F1 Score -avg/total	0,91	0,91	0,91	0,92	0,92	0,87	0,91
Precision-class 1	0,37	0,36	0,37	0,39	0,41	0,27	0,36
Recall-class 1	0,89	0,86	0,92	0,92	0,93	0,97	0,96
F1 Score-class 1	0,53	0,50	0,53	0,54	0,57	0,43	0,53

ROC curve AUC

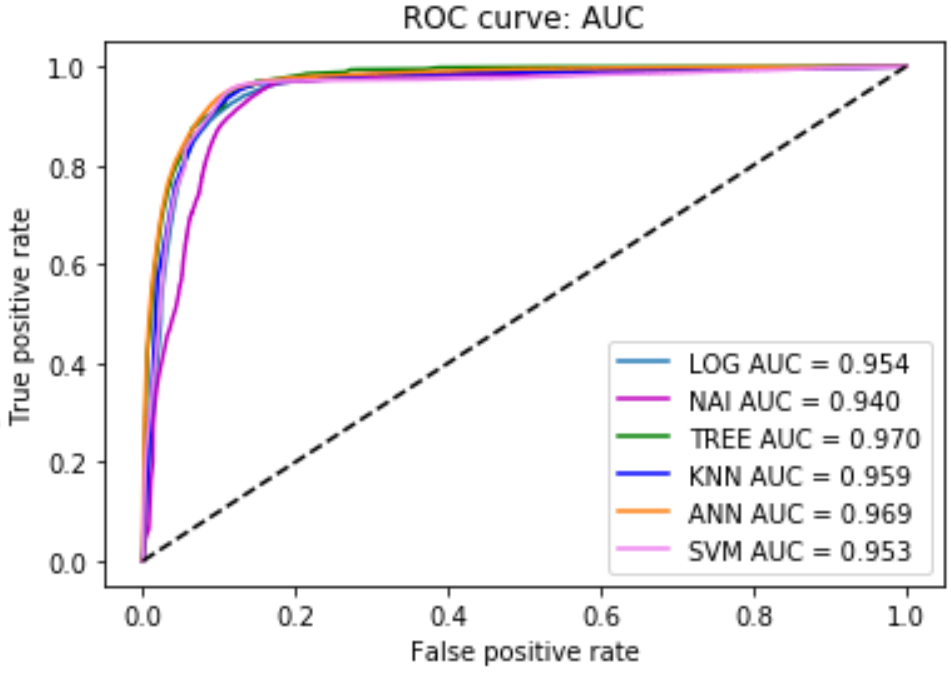

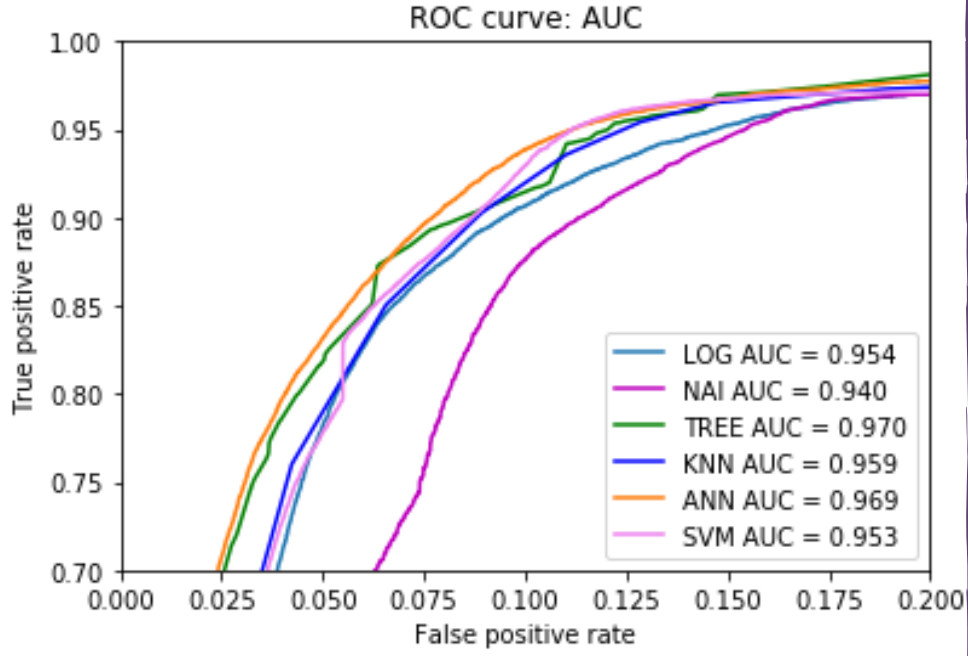

Conclusion and Outlook

#1 Decision Tree

Optimal parameters:
- criterion='entropy',
- min impurity decrease=0.000015,
- max depth=15,
- max leaf nodes=None,
- min samples leaf=10,
- min samples split=20

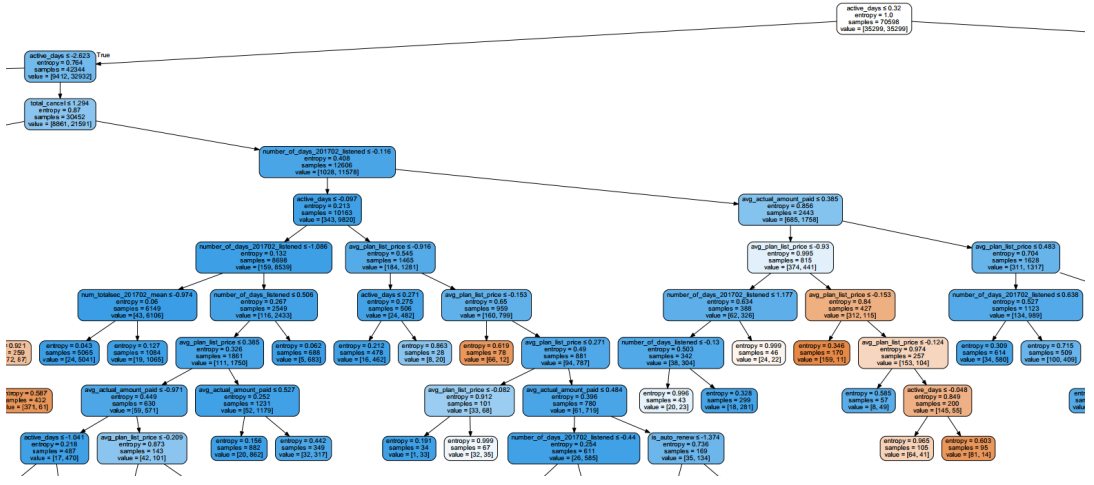

Actual	Predicted: 0	Predicted: 1
0	219,970	25,580
1	1436	15,857

Discussion

Drawbacks:
- other balancing approaches: oversampling
- scale out to the whole timeframe (3 years)

Outlook:
- hybrid model*

*Lee, Jae, and Jin Lee. "Customer churn prediction by hybrid model." Advanced Data Mining and Applications (2006): 959-966.

Thank You !

Presented by
Team 9:

Simona Doneva
Na Gong
Bengi Koseouglu
Siying Liu
Liam Pang

Questions?


YOU ARE A DATA SCIENTIST
YOU ARE A DATA SCIENTIST
YOU ARE A DATA SCIENTIST
YOU ALL ARE DATA SCIENTISTS!

Decision Tree on train data before standardization and balancing

Optimal parameters:
- criterion='entropy',
- min impurity decrease=0.000015,
- max depth=15,
- max leaf nodes=None,
- min samples leaf=10,
- min samples split=20

		Predicted	
		0	1
Actual	0	243,240	2310
	1	8212	9081

Top Features

- active days
- total churn
- total cancel
- actual payment
- registration method
- total_sec_listened
- total_days_listened

....